




Figure 1: **PRISM decomposition of the risk gap.** The cross-entropy risk gap $|\mathcal{R}_T - \mathcal{R}_P|$ between target T and proxy P is bounded by a feature alignment error δ and a head discrepancy γ (Thm. ??). Under Procrustes alignment $W^* = \arg \min_{W \in \mathcal{O}(d)} \|Z_T - Z_P W\|_F$, δ decomposes exactly into *scale mismatch* $(\Delta\rho)^2 = (\rho_T - \rho_P)^2$ and *shape mismatch* $2\rho_T\rho_P(1-\Omega)$